Exploratory Data Analysis: Golub Case Study

Chapter 1, Lab 3

OpenIntro Biostatistics

This lab presents the details of how to conduct the analysis discussed in Section 1.7.2 of *OpenIntro Biostatistics*. We recommend that a reader carefully review the material in the text prior to working through the lab, as the lab will focus on demonstrating techniques in R rather than reviewing the conceptual strategy behind the analysis.

Background information

The 1999 Golub leukemia study represents one of the earliest applications of microarray technology for diagnostic purposes. At the time of the Golub study, no single diagnostic test was sufficient for distinguishing between acute myeloid leukemia (AML) and acute lymphoblastic leukemia (ALL). To investigate whether gene expression profiling could be a tool for classifying acute leukemia type, Golub and co-authors used Affymetrix DNA microarrays to measure the expression level of 7,129 genes from children known to have either AML or ALL.

The goal of the study was to develop a procedure for distinguishing between AML and ALL based only on the gene expression levels of a patient. There are two major issues to be addressed:

1. *Which genes are the most informative for making a prediction?* If a gene is differentially expressed between individuals with AML versus ALL, then measuring the expression level of that gene may be informative for diagnosing leukemia type. For example, if a gene tends to be highly expressed in AML individuals, but only expressed at low levels in ALL individuals, it is more likely to be a reliable predictor of leukemia type than a gene that is expressed at similar levels in both AML and ALL patients.

2. *How can leukemia type be predicted from expression data?* Suppose that a patient's expression profile is measured for a group of genes. In an ideal scenario, all the genes measured would express AML-like expression, or ALL-like expression, making a prediction obvious. In reality, however, a patient's expression profile will not follow an idealized pattern. Some of the genes may have expression levels more typical of AML, while others may suggest ALL. It is necessary to clearly define a strategy for translating raw expression data into a prediction of leukemia type.

All datasets used in this lab are available from the `oibiostat` package. Phenotypic and expression data have been collected for 72 patients. The expression data from the 62 patients in `golub.train` will be used to identify informative genes for making a prediction. The prediction strategy will then be tested on the remaining 10 patients in `golub.test`.

Identifying informative genes

The discussion in the text begins by illustrating concepts using a simplified version of the dataset (golub.small) that contains only data from the 10 patients and 10 genes. Here, instead of starting with golub.small, we will examine a random sample of 100 genes for all patients in golub.train. The methods from the initial analysis can then be applied to the data from all 7,129 genes.

1. Run the following code to load golub.train and create gene.matrix, which contains only the expression data and not the phenotype information in the first 6 columns.

```
#load the data
library(oibiostat)
data(golub.train)

gene.matrix = as.matrix(golub.train[ ,-(1:6)])
```

By using the - in front of the column numbers, the matrix notation specifies that columns 1 through 6 should *not* be included. The same matrix could be created by specifying that columns 7 through 7,135 should be included, with [,7:7135].

2. Draw a random sample of 100 genes from the dataset.

```
#create a vector of integers from 1 to the total number of genes
gene.columns = 1:ncol(gene.matrix)

#set the seed for a pseudo-random sample
set.seed(2401)

#sample 100 numbers from gene.columns, without replacement
gene.index.set = sample(gene.columns, size = 100, replace = FALSE)

#creeate a matrix with expression data from the rows specified by gene.index.set
gene.matrix.sample = gene.matrix[ ,gene.index.set]
```

a) What are the first five values of gene.index.set? How were the numbers in gene.index.set chosen?

b) Why is it important to sample without replacement?

c) View gene.matrix.sample; what does it contain? How is gene.matrix.sample related to gene.index.set?

2

d) What are the first five gene names of the 100 genes sampled?

e) Plot a histogram showing the distribution of the expression levels of the second gene across patients. Describe the distribution.

3. Create a logical variable, leuk.type, that has value 1 for AML and value 0 for anything that is not AML (i.e., allT and allB).

```
#create logical variable
leuk.type = (golub.train$cancer == "aml")

#view table of leukemia types
table(leuk.type)
```

a) When creating the logical variable, why not write "allT" or "allB" instead of "aml"?

b) How many patients have AML? How many have ALL?

4. Summarize the data separately for AML patients and for ALL patients.

a) The following code calculates the mean expression level for each sampled gene across AML patients, storing it in the variable aml.mean.expression. The apply() function executes a function across a matrix—in this case, the function is mean, and the 2 in the argument indicates that the function should be applied on each column (replacing the 2 with a 1 would result in the mean being calculated across the rows).

```
#calculate mean expression level for each sampled gene across AML patients
aml.mean.expression = apply(gene.matrix.sample[leuk.type == TRUE, ], 2, mean)
```

Run the code to create aml.mean.expression, then create all.mean.expression, a vector containing the mean expression levels for each gene in ALL patients.

b) Explain the logic behind the code to generate aml.mean.expression and all.mean.expression. In other words, what do the separate components instruct R to do?

c) View the contents of `aml.mean.expression`. What is the average expression level of the first sampled gene across AML patients?

5. For each gene, compare the mean expression value among AML patients to the mean among ALL patients; calculate the differences in mean expression levels between AML and ALL patients.

```
#calculate the differences
diff.mean.expression = (aml.mean.expression - all.mean.expression)

#view list as a matrix
diff.mean.expression.matrix = as.matrix(diff.mean.expression)
diff.mean.expression.matrix
```

a) What is the difference in mean expression level between AML and ALL for the first gene on the list; on average, is this gene more highly expressed in AML patients or ALL patients? Does it seem like this gene could be a good predictor of leukemia type? Why or why not?

b) Using numerical and graphical summaries, describe the distribution of differences in mean expression levels.

6. Identify the outliers. Run the following code to set up the definition of outliers as specified in Chapter 1 of *OpenIntro Biostatistics*:

```
#define 3rd and 1st quartiles
quart.3 = quantile(diff.mean.expression.matrix[ ,1], 0.75, na.rm = TRUE)
quart.1 = quantile(diff.mean.expression.matrix[ ,1], 0.25, na.rm = TRUE)

#define interquartile range
iqr = quart.3 - quart.1

#define upper and lower bound for outliers
lb.outlier = quart.1 - 1.5*iqr
ub.outlier = quart.3 + 1.5*iqr
```

The following code creates a list of the large outliers, genes with expression differences larger than ub.outlier:

```
#creates list of large outliers
which.large.out = diff.mean.expression > ub.outlier
large.out = as.matrix(diff.mean.expression.matrix[which.large.out, ])
large.out

#creates ordered list of large outliers, from largest to smallest
order.large.out = order(large.out[ ,1], decreasing = TRUE) #assigns ordering to rows
ordered.large.out = as.matrix(large.out[order.large.out, ]) #sorts outlier list
ordered.large.out
```

a) What are the upper and lower outlier bounds?

b) How many large outliers are present in the sample?

c) How many rows and columns does large.out have? Explain why.

d) View order.large.out. What do these numbers represent?

e) Modify the code to find small outliers. How many small outliers are present in the sample?

f) Which gene has the largest positive difference in mean expression between AML and ALL samples? Which gene has the largest negative difference in mean expression between AML and ALL samples?

g) In a research setting, it can also be useful to inspect the entire list and examine genes that are close to the outlier cutoff. Run the following code to order the entire list of expression differences in decreasing order:

```
order.decreasing = order(diff.mean.expression.matrix[ ,1], decreasing = TRUE)
ordered.outliers = as.matrix(diff.mean.expression.matrix[order.decreasing, ])
```

Which gene just missed the cutoff to qualify as a large outlier? Which gene is closest to the cutoff for qualifying as a small outlier?

7. The genes previously identified as outliers are only outliers of the specific 100 genes chosen in the sample. From the complete set of data in golub.train, identify the five largest outliers and five smallest outliers out of all 7,129 genes. (*Hint*: this can be done with only a few modifications to the code run for the initial analysis.)

Predicting leukemia type

Figure 1 illustrates the main ideas behind the strategy developed by the Golub team to predict leukemia type from expression data. The vertical orange bars represent the gene expression levels of a patient for each gene, relative to the mean expression for AML patients and ALL patients from the training dataset (vertical blue bars). A gene will "vote" for either AML or ALL, depending on whether the patient's expression level is closer to μ_{AML} or μ_{ALL}. In the example shown, three of the genes are considered to have ALL-like expression, versus the other two that are more AML-like. The votes are also weighted to account for how far an observation is from the midpoint between the two means (horizontal dotted blue line); i.e., the length of the dotted line shows the deviation from the midpoint. For example, the observed expression value for gene 2 is not as strong an indicator of ALL as the expression value for gene 1. The magnitude of the deviations $(v_1, v_2, \ldots)$ are summed to obtain V_{AML} and V_{ALL}, and a higher value indicates a prediction of either AML or ALL, respectively.


Figure 1: Schematic of the prediction strategy used by the Golub team, reproduced with modifications from Fig. 1B of the original paper.

The published analysis chose to use 50 informative genes. For simplicity, a smaller number of genes will be used in our version of the analysis. As identified in the previous section, 10 genes will be used as predictors—the 5 largest outliers and 5 smallest outliers for the difference in mean expression between AML and ALL. The expression levels for these 10 genes will be used to make predictions of leukemia type for the patients in `golub.test`.

Steps in the analysis:

- Calculate μ_{AML} and μ_{ALL} for each of the predictor genes.
- Determine the vote direction for each gene; is its expression more AML-like or ALL-like?
- Calculate $v_1, \ldots, v_{10}$, the magnitude of the deviations from the midpoint between the two means..
- Calculate V_{AML} and V_{ALL} for each patient, and determine which leukemia type is predicted.
- Assess prediction accuracy, comparing the prediction to the actual leukemia status of each patient; recall that leukemia status is known for all study patients.

8. Run the following code to identify the column numbers in `golub.test` and `golub.train` corresponding to the 10 predictor genes identified in Question 7. Using `which()` returns the column numbers that correspond to the 10 predictor genes. from `gene.matrix`, the matrix

without the 6 columns of phenotypic data. Adding 6 results in the corresponding column numbers in golub.test and golub.train.

```
#column numbers for large outliers (gene.matrix)
predictors.large = which(which.large.out)[order.large.out[1:5]]

#column numbers for small outliers (gene.matrix)
predictors.small = which(which.small.out)[order.small.out[1:5]]

#combine column numbers for large and small outliers (gene.matrix)
predictors = c(predictors.large, predictors.small)

#predictor column numbers of golub.test and golub.train
predictor.cols = predictors + 6
```

9. Calculate μ_{AML} and μ_{ALL} for each of the predictor genes. Store the results in predictor.means.aml and predictor.means.all. What are μ_{AML} and μ_{ALL} for the first predictor gene?

10. Determine the vote direction for each gene. Let 0 represent a vote for AML and 1 represent a vote for ALL.

```
#load golub.test
data("golub.test")

#create matrix with expression data for predictor genes
test.gene.matrix = golub.test[, predictor.cols]

#create empty matrix to store vote directions
num.genes = ncol(test.gene.matrix); num.patients = nrow(test.gene.matrix)
votes = matrix(nrow = num.patients, ncol = num.genes)

#calculate vote directions
dist.from.aml = vector("numeric", num.genes)
dist.from.all = vector("numeric", num.genes)

for(i in 1:dim(votes)[1]){
  for(j in 1:dim(votes)[2]) {

dist.from.aml[i] = abs(test.gene.matrix[i, j] - predictor.means.aml[j])
dist.from.all[i] = abs(test.gene.matrix[i, j] - predictor.means.all[j])

if(dist.from.aml[i] <= dist.from.all[i]){
  votes[i, j] = 0  #assigns 0 if dist. from AML mean <= dist. from ALL mean
} else { votes[i, j] = 1  #assigns 1 if otherwise
}


  }
```

```
}
```

The above code uses `for()` loops and a conditional `if()` statement to assign vote directions for each of the 10 predictor genes, for each patient. A more formal introduction to loops and conditional statements will be provided in the next chapter; for now, focus on understanding the logic of the code, rather than the precise syntax.

a) The loop stores the results in a matrix called votes. What are the dimensions of the matrix? Does a single row contain the votes for a single patient, or the votes for a single gene? (*Hint*: refer to the syntax used to create the empty votes matrix.)

b) Inside the loop, the two vectors `dist.from.aml` and `dist.from.all` store the distance between an expression value and the AML or ALL mean, respectively; the distance is calculated as the absolute value of the expression value in a specific cell minus a predictor mean (either AML or ALL). The loop extends in two directions, where i ranges from 1:10 and j ranges from 1:10, since the dimensions of the votes matrix are 10 (rows) by 10 (columns).

  i. In the first step of the loop, i = 1 and j = 1. Based on what you know about bracket notation, describe what is being calculated and stored in `dist.from.aml` and `dist.from.all` for these values of i and j.

  ii. What is being calculated and stored when i = 2 and j = 1?

  iii. What is being calculated and stored when i = 1 and j = 2?

  iv. Based on the previous answers, why does it make sense for the loop to calculate distance based on `predictor.means.aml[j]` and `predictor.means.all[j]` instead of `predictor.means.aml[i]` and `predictor.means.all[i]`?

c) View the first row of the matrix votes. For which predictor genes does this patient have AML-like expression, and for which does this patient have ALL-like expression?

11. Calculate $v_1, \ldots, v_{10}$, the magnitude of the deviations from the midpoint between the two means. The following code stores the magnitudes of the deviations in a matrix named `deviation.magnitude`.

```
deviation.magnitude = matrix(nrow = num.patients, ncol = num.genes)

for(i in 1:dim(deviation.magnitude)[1]){
  for(j in 1:dim(deviation.magnitude)[2]) {

midpoint = (predictor.means.aml - predictor.means.all)/2

deviation.magnitude[i,j] = abs(test.gene.matrix[i, j] - midpoint[j])

  }
}

#adds predictor gene probe names from test.gene.matrix
colnames(deviation.magnitude) <- colnames(test.gene.matrix)
```

a) View `deviation.magnitude`. What is the deviation from the midpoint at the M19507_at probe for patient 1? For which patient is the deviation at this gene probe the largest?

b) For patient 1, which values should be summed to calculate V_{AML}? (*Hint*: Refer to Question 10, part c.)

12. Calculate V_{AML} and V_{ALL} for each patient, and determine which leukemia type is predicted.

```
#sum the votes for AML and ALL
V.aml = vector("numeric", num.patients)
V.all = vector("numeric", num.patients)

for(i in 1:num.patients){

  V.aml[i] = sum(deviation.magnitude[i, which(votes[i,] == 0)])
  V.all[i] = sum(deviation.magnitude[i, which(votes[i,] == 1)])

}
```

```
#determine the predicted leukemia type
predicted.leuk.type = vector("numeric", num.patients)

for(i in 1:num.patients){
  if (V.aml[i] > V.all[i]){predicted.leuk.type[i] = "aml"}
  if (V.aml[i] < V.all[i]){predicted.leuk.type[i] = "all"}
  if (V.aml[i] == V.all[i]){predicted.leuk.type[i] = "tie"}
}

predicted.leuk.type
```

a) Briefly describe how the `for()` loop to create `V.aml` and `V.all` works.

b) Briefly describe how the `for()` loop and `if` statements to create `predicted.leuk.type` work.

c) What are the predicted leukemia types for the 10 patients in `golub.test`?

13. Assess the prediction accuracy, comparing the prediction to the actual leukemia status of each patient in `golub.test`. How well do the predictions match patient leukemia status?

```
#compare predictions to actual leukemia type

#option 1: visual comparison
golub.test$cancer
predicted.leuk.type


#option 2: use logical variables
actual.leuk.type.ind = (golub.test$cancer == "aml")
predicted.leuk.type.ind = (predicted.leuk.type == "aml")

comparison.results = (actual.leuk.type.ind == predicted.leuk.type.ind)
table(comparison.results)
```