

# Lab Notes

## *Chapter 5*

### *OpenIntro Biostatistics*

## Overview

1. Two-Sample Tests
  - *OI Biostat* Sections 5.2 - 5.3
2. Statistical Power
  - *OI Biostat* Section 5.4
3. Analysis of Variance (ANOVA)
  - *OI Biostat* Section 5.5
4. Multiple Testing
  - Extension to *OI Biostat* Section 5.5
5. A Closer Look at the  $P$ -Value
  - Extension to *OI Biostat* Section 5.6

Lab 1 introduces hypothesis testing in the two-sample context, discussing the two-sample  $t$ -test for paired data and independent group data.

Lab 2 discusses the control of Type I and Type II error and explores the factors influencing the power of a statistical test via simulation.

Lab 3 introduces the analysis of variance procedure for comparing the means of several groups.

Lab 4 examines the multiple testing problem and concept of experiment-wise error in the context of the Golub leukemia data.

Lab 5 integrates the ideas of conditional probability and hypothesis testing to present a broader understanding of  $p$ -values, Type I error, and statistical power in a research context.

## Lab 1: Two-Sample Tests

### Hypothesis Testing with `t.test()`, cont.

The `t.test()` function has the following generic structure:

```
t.test(x, y, alternative = "two.sided", mu = 0, conf.level = 0.95, paired = FALSE)
```

where `x` and `y` are numeric vectors of data values, `alternative` specifies the form of the alternative hypothesis, `mu` is  $\mu_1 - \mu_2$  (in the paired context,  $\delta_0$ ), and `conf.level` refers to the confidence level. The argument for `alternative` can be either "two.sided" ( $H_A : \mu_1 \neq \mu_2$ ), "less" ( $H_A : \mu_1 < \mu_2$ ), or "greater" ( $H_A : \mu_1 > \mu_2$ ). By default, confidence level is set to 95%, and a two-sided alternative is tested with the independent group test.

To conduct a test on data contained in variable `y` that is grouped by the variable `x`, use the tilde syntax:

```
t.test(y ~ x, ...)
```

The following example shows a hypothesis test for mean standing height in centimeters in the artificial NHANES population, using a random sample of 135 adults. The null hypothesis is that the population mean height for females is equal to the population mean height for males. A one-sided alternative is tested against the null; the output includes the  $t$ -statistic, degrees of freedom,  $p$ -value, 90% confidence interval, and the sample means of both groups.

```
#load the data
library(oibiostat)
data("nhanes.samp.adult")

#conduct test
t.test(nhanes.samp.adult$Height ~ nhanes.samp.adult$Gender, alternative = "less",
       conf.level = 0.90)
```

```
##
##  Welch Two Sample t-test
##
## data:  nhanes.samp.adult$Height by nhanes.samp.adult$Gender
## t = -10.777, df = 132.95, p-value < 2.2e-16
## alternative hypothesis: true difference in means is less than 0
## 90 percent confidence interval:
##      -Inf -11.91326
## sample estimates:
## mean in group female   mean in group male
##           162.9729           176.5031
```

The following example shows two ways to conduct a hypothesis test for the difference in mean maximal swim velocity between swimmers wearing wetsuits versus swimsuits. The data are paired, since each participant completed two trials: one wearing a wetsuit and one wearing a swimsuit. The null hypothesis of no difference of  $H_0 : \delta = 0$  is tested against the two-sided alternative  $H_A : \delta \neq 0$ .

The first method uses the two-sample test syntax, while the second method uses the one-sample test syntax on the vector of velocity differences.

```
#load the data
library(oibiostat)
data("swim")

#two-sample test syntax
t.test(swim$wet.suit.velocity, swim$swim.suit.velocity, alternative = "two.sided",
        paired = TRUE)
```

```
##
## Paired t-test
##
## data: swim$wet.suit.velocity and swim$swim.suit.velocity
## t = 3.7019, df = 11, p-value = 0.00349
## alternative hypothesis: true difference in means is not equal to 0
## 95 percent confidence interval:
##  0.02534062 0.09965938
## sample estimates:
## mean of the differences
##                0.0625
```

```
#one-sample test syntax
t.test(swim$velocity.diff, mu = 0, alternative = "two.sided")
```

```
##
## One Sample t-test
##
## data: swim$velocity.diff
## t = 3.7019, df = 11, p-value = 0.00349
## alternative hypothesis: true mean is not equal to 0
## 95 percent confidence interval:
##  0.02534062 0.09965938
## sample estimates:
## mean of x
##    0.0625
```

## Lab 2: Statistical Power

### Simulating Values from a Distribution

R has built-in functions for drawing random values from a distribution. The function `rnorm()` is used in Lab 2 to draw observations from normal distributions with specified parameter values. For reference, details for sampling values from other distributions are also discussed in this section.

The function `rnorm()` has the generic structure

```
rnorm(n, mean = 0, sd = 1)
```

where `n` is the number of observations sampled. By default, R assumes that mean and standard deviation are 0 and 1, respectively.

The following code shows how to draw 10 values from a normal distribution with mean 100 and standard deviation 5. As with any random sampling, it is necessary to specify a seed with `set.seed()` for the results to be reproducible.

```
#set seed for pseudorandom sampling
set.seed(2018)

#draw values
rnorm(10, mean = 100, sd = 5)
```

```
## [1] 97.88508 92.25061 99.67785 101.35441 108.67642 98.67644 110.49735
## [8] 104.31676 96.94706 103.18528
```

The function `rbinom()` has the generic structure

```
rbinom(n, size, prob)
```

where `n` is the number of observations sampled, `size` is the number of trials `n`, and `prob` is the probability of success  $p$ .

The following code shows how to draw 10 values from a binomial distribution with 10 trials and success probability 0.35.

```
rbinom(10, 10, 0.35)
```

```
## [1] 2 4 2 2 5 4 3 6 4 1
```

The function `rpois()` has the generic structure

```
rpois(n, lambda)
```

where `n` is the number of observations sampled and `lambda` is the rate parameter  $\lambda$ .

The following code shows how to draw 10 values from a Poisson distribution with rate parameter  $\lambda = 3$ .

```
rpois(10, 3)
```

```
## [1] 1 5 0 3 3 4 0 1 3 1
```

The function **rgeom()** has the generic structure

```
rgeom(n, prob)
```

where  $n$  is the number of observations sampled and  $prob$  is the probability of success  $p$ .

The following code shows how to draw 10 values from a geometric distribution with probability of success  $p = 0.35$ .

```
rgeom(10, 0.35)
```

```
## [1] 1 0 0 1 4 2 3 1 2 0
```

The function **rnbinom()** has the generic structure

```
rnbinom(n, size, prob)
```

where  $n$  is the number of observations sampled,  $size$  is the number of successes  $r$ , and  $prob$  is the probability of success  $p$ .

The following code shows how to draw 10 values from a negative binomial distribution with number of successes  $r = 4$  and probability of success  $p = 0.8$ .

```
rnbinom(10, 4, 0.8)
```

```
## [1] 0 0 0 0 0 2 1 0 1 1
```

The function **rhyper()** has the generic structure

```
rhyper(nn, m, n, k)
```

where  $nn$  is the number of observations sampled,  $m$  is the total number of successes  $m$ ,  $n$  is the total number of failures  $N - m$ , and  $k$  is the sample size  $n$ .

The following code shows how to draw 10 values from a hypergeometric distribution with total number of successes  $m = 10$ , total number of failures  $N - m = 15$ , and sample size  $n = 8$ .

```
rhyper(10, 10, 15, 8)
```

```
## [1] 3 3 3 4 3 2 2 3 3 4
```

## Power and Sample Size Calculations with **power.t.test()**

The **power.t.test()** function can both compute the power of a one- or two-sample  $t$ -test and determine necessary parameters (e.g., sample size) to obtain a target power. The function has the generic structure

```
power.t.test(n = NULL, delta = NULL, sd = 1, sig.level = 0.05,  
             power = NULL, type, alternative)
```

where  $n$  is the sample size (per group),  $\delta$  is the effect size,  $sd$  is the standard deviation,  $\text{sig.level}$  is the significance level, and  $\text{power}$  is the statistical power. The argument for `type` can be either "one.sample", "two.sample", or "paired", where two-sample implies independent groups. The argument for `alternative` can be either "two.sided" or "one.sided".

Exactly one out of  $n$ ,  $\delta$ ,  $sd$ , or  $\text{sig.level}$  must be entered as NULL; this is the parameter of interest that will be calculated based on the provided information.

The following code shows how to calculate the power for a one-sample test where  $n = 100$ ,  $\Delta = 3$ ,  $\sigma = 12$ ,  $\alpha = 0.05$ , with a two-sided alternative.

```
power.t.test(n = 100, delta = 3, sd = 12, sig.level = 0.05,  
             power = NULL, type = "one.sample", alternative = "two.sided")
```

```
##  
##      One-sample t test power calculation  
##  
##              n = 100  
##            delta = 3  
##              sd = 12  
##      sig.level = 0.05  
##            power = 0.6969757  
##      alternative = two.sided
```

The following code shows how to calculate the sample size for a one-sample test where  $\Delta = 3$ ,  $\sigma = 12$ ,  $\alpha = 0.05$ , and power of 0.70, with a two-sided alternative.

```
power.t.test(n = NULL, delta = 3, sd = 12, sig.level = 0.05,  
             power = 0.70, type = "one.sample", alternative = "two.sided")
```

```
##  
##      One-sample t test power calculation  
##  
##              n = 100.6887  
##            delta = 3  
##              sd = 12  
##      sig.level = 0.05  
##            power = 0.7  
##      alternative = two.sided
```

## Lab 3: Analysis of Variance (ANOVA)

### The `tapply()` Function

The **`tapply()`** function is related to the `apply()` function introduced in Chapter 1. As with `apply()`, `tapply()` allows a specific function to be applied to a matrix; the function can be a pre-defined R function like `mean()` or a user-defined function. The power of `tapply()` is that it allows for a vector to be split into groups, with the function applied to each group.

The function has the generic structure

```
tapply(y, x, FUN)
```

where `y` is the vector of data, `x` is the grouping variable, and `FUN` is the function of interest.

The following code shows how to calculate the mean change in non-dominant arm strength for each genotype group in the FAMuSS data.

```
#load the data
library(oibiostat)
data("famuss")

tapply(famuss$ndrm.ch, famuss$actn3.r577x, mean)
```

```
##          CC          CT          TT
## 48.89422 53.24904 58.08385
```

### Fitting an ANOVA Model

The **`aov()`** function fits an ANOVA model to data; wrapping with the `summary()` function outputs the ANOVA table, which contains the  $F$ -statistic and associated  $p$ -value. The input to `aov()` must be in the form of a formula using the tilde syntax:

```
aov(y ~ x)
```

where `y` is the data vector and `x` is the grouping variable.

The following code shows the summary of the ANOVA model fit for the association of change in non-dominant arm strength by genotype at the *r577x* locus on the *ACTN3* gene.

```
#output summary of anova model
summary(aov(famuss$ndrm.ch ~ famuss$actn3.r577x))

##              Df Sum Sq Mean Sq F value Pr(>F)
## famuss$actn3.r577x    2    7043    3522   3.231 0.0402 *
## Residuals          592  645293    1090
## ---
## Signif. codes:  0 '***' 0.001 '**' 0.01 '*' 0.05 '.' 0.1 ' ' 1
```

## Conducting Pairwise Tests with `pairwise.t.test`

The `pairwise.t.test` function is used to conduct pairwise comparisons with corrections for multiple testing. Note that the input to this function uses different syntax from `aov()`: instead of the tilde, the comma is used to separate the data variable and grouping variable. The generic structure of the function is

```
pairwise.t.test(y, x, p.adj)
```

where `y` is the data vector, `x` is the grouping vector, and `p.adj` can be one of several adjustment choices, such as "none" for no correction and "bonf" for Bonferroni.

The following code shows how to conduct pairwise two-sample *t*-tests between mean change in non-dominant arm strength for each of the genotype groups in the FAMuSS data.

```
#no correction
pairwise.t.test(famuss$ndrm.ch, famuss$actn3.r577x, p.adj = "none")
```

```
##
## Pairwise comparisons using t tests with pooled SD
##
## data: famuss$ndrm.ch and famuss$actn3.r577x
##
##      CC      CT
## CT 0.179 -
## TT 0.011 0.144
##
## P value adjustment method: none
```

```
#Bonferroni correction
pairwise.t.test(famuss$ndrm.ch, famuss$actn3.r577x, p.adj = "bonf")
```

```
##
## Pairwise comparisons using t tests with pooled SD
##
## data: famuss$ndrm.ch and famuss$actn3.r577x
##
##      CC      CT
## CT 0.537 -
## TT 0.034 0.433
##
## P value adjustment method: bonferroni
```

Note that when the Bonferroni correction is applied, R multiplies the *p*-value by *K*, the number of comparisons; thus, the values output from `pairwise.t.test` when `p.adj = "bonf"` should be compared to  $\alpha$ , not  $\alpha^*$ . Comparing an unadjusted *p*-value to  $\alpha/K$  is equivalent to comparing the quantity ( $K \times p$ -value) to  $\alpha$ .



## Lab 4: Multiple Testing

The for loop was introduced in Chapter 2; nested loops were introduced in Chapter 3 in the context of simulating geometric, negative binomial, and hypergeometric random variables. This section specifically discusses nested for loops and the logic behind the simulation code for estimating experiment-wise error in the Golub leukemia dataset.

### Nested for Loops

Understanding a nested for loop requires keeping track of more than two counters (i.e., index variables). In the following basic example, there are two counters: the outer counter,  $k$ , runs from 1 through 4, while the inner counter,  $j$ , runs from 1 through 2.

- For the first iteration,  $k = 1$ . Upon encountering the second loop, R cycles through  $j = 1$  and  $j = 2$ . Thus, there are two values of the product  $k \times j$  for this first iteration:  $1 \times 1 = 1$  and  $1 \times 2 = 2$ .
- For the fourth iteration,  $k = 4$ . The two values of the product  $k \times j$  are then  $4 \times 1 = 4$  and  $4 \times 2 = 8$ .

```
for(k in 1:4){  
  
  for(j in 1:2){  
    print(k*j)  
  }  
  
}
```

```
## [1] 1  
## [1] 2  
## [1] 2  
## [1] 4  
## [1] 3  
## [1] 6  
## [1] 4  
## [1] 8
```

Question 2 of the lab refers to a simulation for estimating experiment-wise error rate when two independent one-sample hypothesis tests are conducted. The approach shown in the lab is to create two separate vectors of observations. While this approach is straightforward, it is impractical for a large number of tests.

The following code demonstrates a more flexible approach that hinges on using nested for loops. When the number of tests is specified as a parameter, the simulation can simply be re-run to model experiment-wise error for any number of tests.

- The outer loop, with index variable  $k$ , runs from 1 to the specified number of iterations. The inner loop, with index variable  $j$ , runs from 1 to the specified number of tests.
- Each time the outer loop runs, a set of data (i.e., observations in samples to be tested) is

generated. The matrix `obs.matrix` has number of columns equivalent to `num.tests` and number of rows equivalent to `num.obs`. It is populated with  $num.tests \times num.obs = 100 \times 100$  draws from a standard normal distribution. This is a more efficient way to generate the simulated data than running `rnorm()` 100 times and creating 100 vectors.

- The inner loop proceeds through each column of `obs.matrix`, conducting a  $t$ -test on the values in column  $j$  and storing the  $p$ -value as the  $j^{th}$  entry of the vector `p.vals`.
- The last instruction in the outer loop is to record the minimum value in `p.vals` as the  $k^{th}$  entry in the vector `min.p.vals`. Note how the values in `obs.matrix` and `p.vals` are rewritten with each iteration of the outer loop, but not the values in `min.p.vals`.
- The reject vector is also defined more efficiently than in the version of the simulation shown in the lab. If the  $k^{th}$  entry in `min.p.vals` is larger than  $\alpha$ , then the  $k^{th}$  iteration represents one instance of experiment-wise error occurring. This logic was discussed in Question 7.
- The result of the simulation agrees closely with the algebraic solution from Question 3 of the lab. From simulation, the estimated experiment-wise error is 0.995; the probability of at least one incorrect rejection in 100 independent tests conducted at  $\alpha = 0.05$  is 0.994.

```
#set parameters
num.tests = 100
num.obs = 100
num.iterations = 1000

alpha = 0.05

#set seed
set.seed(2018)

#create empty lists
p.vals = vector("numeric", num.tests)
min.p.vals = vector("numeric", num.iterations)

#run simulation
for(k in 1:num.iterations){

  obs.matrix = matrix(rnorm(num.tests*num.obs),
                      nrow = num.obs, ncol = num.tests)

  for(j in 1:num.tests){

    p.vals[j] = t.test(obs.matrix[, j], mu = 0)$p.val

  }

  min.p.vals[k] = min(p.vals)

}
```

```
#view results  
reject = (min.p.vals <= alpha)  
table(reject)
```

```
## reject  
## FALSE TRUE  
##      5  995
```

## Lab 5: A Closer Look at the $P$ -Value

The simulation for estimating posterior probabilities follows the same logical structure as used when simulating populations based on known conditional probabilities (as discussed in Chapter 2), such as in a diagnostic testing context.

In fact, the language of diagnostic testing can be used to frame the ideas of prior and posterior probabilities in hypothesis testing. Suppose that the alternative hypothesis represents presence of a disease condition (i.e., think of the event  $H_A$  is true as the event of having a disease,  $D$ ), and that to reject the null hypothesis is to obtain a positive test result (i.e., think of rejecting  $H_0$  as the event  $T^+$ ). The posterior probability of rejecting  $H_0$  when  $H_A$ , given that  $H_0$  was rejected, is analogous to the positive predictive value  $P(D|T^+)$ .

The loop used in the lab is reproduced here for reference:

- Prior to running the loop, the state of nature (whether  $H_0$  or  $H_A$  is true) was randomly assigned using `sample()` and stored in the vector `hypothesis`. This is akin to randomly assigning which individuals in the population have the disease and which do not.
- The first if statement has the condition `hypothesis[k] == "null"`. If the  $k^{th}$  element of `hypothesis` is null, then the data are simulated under the null hypothesis where the control and treatment means are equal. This is comparable to simulating the number of positive tests obtained when an individual does not have the disease.
- The second if statement has the condition `hypothesis[k] == "alternative"`. If the  $k^{th}$  element of `hypothesis` is alternative, then the data are simulated under the alternative hypothesis where the control and treatment means are not equal. This is comparable to simulating the number of positive tests obtained when an individual has the disease.

```
for(k in 1:num.iterations){  
  
  if(hypothesis[k] == "null"){  
  
    control = rnorm(n.control, mean = control.mean.null, sd = control.sigma)  
    treatment = rnorm(n.treatment, mean = treatment.mean.null, sd = treatment.sigma)  
  
    p.vals[k] = t.test(control, treatment, alternative = "two.sided",  
                       mu = 0, conf.level = 1 - alpha)$p.val  
  
  }  
  
  if(hypothesis[k] == "alternative"){  
  
    control = rnorm(n = n.control, mean = control.mean.alternative, sd = control.sigma)  
    treatment = rnorm(n = n.treatment, mean = treatment.mean.alternative, sd = treatment.sigma)  
  
    p.vals[k] = t.test(control, treatment, alternative = "two.sided",  
                       mu = 0, conf.level = 1 - alpha)$p.val  
  
  }  
  
}
```