

Lab Notes

Chapter 7

OpenIntro Biostatistics

Overview

1. Introduction to Multiple Regression
 - *OI Biostat* Sections 7.1 and 7.2
2. Evaluating Model Fit
 - *OI Biostat* Section 7.3
3. Categorical Predictors with Several Levels and Inference in Regression
 - *OI Biostat* Sections 7.4 - 7.6 and 7.9
4. Interaction
 - *OI Biostat* Section 7.7
5. Model Selection for Explanatory Models
 - *OI Biostat* Section 7.8

Lab 1 introduces the multiple regression model in the context of estimating an association between a response variable and primary predictor of interest while adjusting for possible confounding variables.

Lab 2 discusses the use of residual plots to check assumptions for multiple regression and introduces adjusted R^2 .

Lab 3 extends on the topics introduced in Chapter 6, Lab 4 by discussing categorical predictors with more than two levels and generalizing inference in regression to the setting where there are several slope parameters.

Lab 4 introduces the concept of a statistical interaction, specifically in the case of an interaction between a categorical variable and a numerical variable.

Lab 5 discusses explanatory modeling, in which the goal is to construct a model that explains the observed variation in the response variable. This is an application of multiple regression distinct from that presented in Lab 1.

Lab 1: Introduction to Multiple Regression

Working with Several Predictors

The `lm()` function is used to fit linear models. It has the following generic structure:

```
lm(y ~ x1 + x2, data)
```

where the first argument specifies the variables used in the model; in this example, the model regresses a response variable y against two explanatory variables x_1 and x_2 . Additional predictor variables can be added to the model formula with the `+` symbol.

The following example shows fitting a linear model that predicts BMI from age (in years) and gender using data from `nhanes.samp.adult.500`, a sample of individuals 21 years of age or older from the NHANES data.

```
#load the data
library(oibistat)
data("nhanes.samp.adult.500")

#fitting linear model
lm(BMI ~ Age + Gender, data = nhanes.samp.adult.500)

##
## Call:
## lm(formula = BMI ~ Age + Gender, data = nhanes.samp.adult.500)
##
## Coefficients:
## (Intercept)      Age  Gendermale
##    28.8087      0.0206     -0.9571
```

Letting R do the Work: Predicted Values

The `predict()` function can be used to evaluate the regression equation for specific x -values, or in other words, to calculate \hat{y} values for values of x that were not necessarily observed. To use `predict()` in this way, specify the x -values according to the following generic syntax:

```
predict(object, newdata = data.frame( ))
```

where `object` is the name of the fitted model, and the name of the predictor variable and value at which to evaluate the equation are specified within `newdata = data.frame()`.

In a model with several variables, values for all variables in the model must be specified to calculate a prediction.

The following example shows calculating \widehat{BMI} for a male individual 60 years of age using the model regressing BMI on age and gender in `nhanes.samp.adult.500`, then checking the result by explicitly solving the regression equation.

```
#BMI ~ Age + Gender in nhanes.samp.adult.500
model.BMIvsAgeGender = lm(BMI ~ Age + Gender, data = nhanes.samp.adult.500)
predict(model.BMIvsAgeGender, newdata = data.frame(Age = 60, Gender = "male"))
```

```
##      1
## 29.09

#confirm answer from solving  $28.81 + 0.02(60) - 0.95(1)$ 
coef(model.BMIvsAgeGender)[1] + coef(model.BMIvsAgeGender)[2]*60 +
  coef(model.BMIvsAgeGender)[3]*1

## (Intercept)
##      29.09
```

Plotting Points According to a Condition

The following plot was generated by using `plot()` to plot a set of points with the associated plot axes; in this example, the cases representing statin non-users were plotted first. Afterwards, the `points()` function was used to overlay a set of points onto the existing plot; these points in red represent statin users. The `points()` function takes the same arguments as `plot()`.

Refer to the notes for Lab 3 for details about using COL to specify colors.


```
#load the data
library(openintro)
library(oibiostat)
data(prevend.samp)
data(COL)

#create statin.use logical
statin.use = (prevend.samp$Statin == 1)

#plot blue points
plot(prevend.samp$Age[statin.use == FALSE], prevend.samp$RFFT[statin.use == FALSE],
     pch = 21, bg = COL[1, 3], col = COL[1], cex = 1.3,
     xlab = "Age (yrs)", ylab = "RFFT Score",
     main = "RFFT Score versus Age (n = 500)")

#plot red points
points(prevend.samp$Age[statin.use == TRUE], prevend.samp$RFFT[statin.use == TRUE],
       pch = 21, bg = COL[4, 3], col = COL[4], cex = 1.3)

#draw vertical lines
abline(v = 40, lty = 2)
abline(v = 50, lty = 2)
abline(v = 60, lty = 2)
abline(v = 70, lty = 2)
abline(v = 80, lty = 2)
```



Lab 2: Evaluating Model Fit

Extracting Adjusted R^2 from a Model Fit

The adjusted R^2 of a model fit can be extracted directly from a model summary with the use of the `$` syntax. For example, the following syntax prints the adjusted R^2 from the model regressing BMI on age and gender in `nhanes.samp.adult.500`.

```
#print adjusted R-squared value
summary(model.BMIvsAgeGender)$adj.r.squared

## [1] 0.0032735
```

Plotting Points According to Several Conditions

The following plots were generated by defining logical variables for ethnicity and logical variables for age. The logical operator `&` was used to succinctly specify which data values should appear in each plot, and specify the display color.

```
#create hispanic and white.not.hisp logicals
hispanic = (dds.subset$ethnicity == "Hispanic")
white.not.hisp = (dds.subset$ethnicity == "White not Hispanic")

#create age logicals
younger = (dds.subset$age < 21)
older = (dds.subset$age >= 21)

par(mfrow = c(1, 2))


#plot blue points, white not hispanic
plot(expenditures[white.not.hisp & younger] ~ age[white.not.hisp & younger],
     data = dds.subset,
     pch = 21, bg = COL[1, 4], col = COL[1], cex = 0.8,
     xlab = "Age (yrs)", ylab = "Expenditures",
     main = "Expenditures vs Age in DDS (0 - 21)")

#plot red points, hispanic
points(expenditures[hispanic & younger] ~ age[hispanic & younger],
       data = dds.subset, pch = 21, bg = COL[4, 4], col = COL[4],
       cex = 0.8)


#plot blue points, white not hispanic
plot(expenditures[white.not.hisp & older] ~ age[white.not.hisp & older],
     data = dds.subset,
     pch = 21, bg = COL[1, 4], col = COL[1], cex = 0.8,
     xlab = "Age (yrs)", ylab = "Expenditures",
     main = "Expenditures vs Age in DDS (21+)")
```

```
#plot red points, hispanic
points(expenditures[hispanic & older] ~ age[hispanic & older],
       data = dds.subset, pch = 21, bg = COL[4, 4], col = COL[4],
       cex = 0.8)
```

Expenditures vs Age in DDS (0 – 21)



Expenditures vs Age in DDS (21+)



Lab 3: Categorical Predictors with Several Levels and Inference in Regression

Color Palettes with RColorBrewer


The RColorBrewer package is a helpful tool for selecting colors. The package provides three types of palettes: sequential palettes, diverging palettes, and qualitative palettes.

- Sequential palettes are ideal for ordered data that progress across a range from low to high.
- Diverging palettes place equal emphasis on extremes at both ends of the data range.
- Qualitative palettes are useful for unordered categorical data.

The `brewer.pal()` function has the generic structure

```
brewer.pal(n, name)
```

where the first argument specifies the number of colors and the second specifies the name of the palette. The possible palette choices are shown below.



The following example shows a figure created using the GnBu palette, one of the sequential palettes. A color scheme progressing from a light shade to a dark shade highlights the ordinal nature of educational level and adds interpretive value to the visual.


```
#load RColorBrewer package
library(RColorBrewer)

#load data
data("prevend.samp")

#convert Education to a factor
prevend.samp$Education = factor(prevend.samp$Education,
                                levels = c(0, 1, 2, 3),
                                labels = c("Primary", "LowerSecond",
                                           "HigherSecond", "Univ"))
```

```
#create plot
plot(RFFT ~ Education, data = prevend.samp,
     xlab = "Educational Level", ylab = "RFFT Score",
     main = "RFFT by Education in PREVEND (n = 500)",
     names = c("Primary", "LowSec", "HighSec", "Univ"),
     col = brewer.pal(4, "GnBu"))
```


RFFT by Education in PREVEND (n = 500)



Using Colors from OpenIntro

The core colors used for the *OpenIntro Statistics* textbook are accessible as the COL dataset in the openintro package.

Each of the seven colors shown below have four levels of transparency, which can be useful when making scatterplots that display a large number of observations.





The colors can be accessed with `COL[]`. A single number within the square brackets identifies each color. For example, `COL[1]` corresponds to blue. A second number within the square brackets can be optionally specified to change the transparency, where 4 is the most transparent and 1 is opaque.

For example, the plot below on the left is made using the *OpenIntro* color package. The dots are outlined in an opaque blue (`col = COL[1]`) while the fill color is in a transparent blue (`col = COL[1, 4]`).

```
#simulate data
x <- rnorm(250)
error <- rnorm(250, 0, 50)
y <- 1.3*x^2 + 25*x + error

par(mfrow = c(1, 2))

#plot with openintro colors
plot(y ~ x,
      pch = 21, col = COL[1], bg = COL[1, 4])
```



The plot on the right is made using `rgb()`, a function for specifying color using the RGB color codes. The `rgb()` function has the generic structure

```
rgb(red, green, blue, max, alpha)
```

where red, blue, green, and alpha can be entered as integers if max is set to 255. The setting alpha is used to specify transparency; a value closer to 0 is more transparent. By default, a color specified will be opaque.

```
#plot with rgb( )  
plot(y ~ x, pch = 21,  
      col = rgb(52, 139, 221, max = 255),  
      bg = rgb(62, 160, 221, max = 255, alpha = 100))
```

Lab 4: Interaction

To fit a model with an interaction term and its main effects, use the `*` symbol in `lm()`. The following example shows fitting a linear model that predicts BMI from age, gender, and the interaction between age and gender in `nhanes.samp.adult.500`.

```
lm(BMI ~ Age*Gender, data = nhanes.samp.adult.500)
```

```
##
## Call:
## lm(formula = BMI ~ Age * Gender, data = nhanes.samp.adult.500)
##
## Coefficients:
##      (Intercept)          Age      Gendermale  Age:Gendermale
##      28.69400         0.02296        -0.69234        -0.00529
```

This is equivalent to explicitly specifying the main effects and the interaction term individually with the `:` symbol.

```
lm(BMI ~ Age + Gender + Age:Gender, data = nhanes.samp.adult.500)
```

```
##
## Call:
## lm(formula = BMI ~ Age + Gender + Age:Gender, data = nhanes.samp.adult.500)
##
## Coefficients:
##      (Intercept)          Age      Gendermale  Age:Gendermale
##      28.69400         0.02296        -0.69234        -0.00529
```

It is also possible to use `()^2`, which indicates to R that all main effects and “second-order” interaction terms should be fit. Interaction terms involving more than two terms are not discussed in this course, but this syntax could be useful in such a setting.

```
lm(BMI ~ (Age + Gender)^2, data = nhanes.samp.adult.500)
```

```
##
## Call:
## lm(formula = BMI ~ (Age + Gender)^2, data = nhanes.samp.adult.500)
##
## Coefficients:
##      (Intercept)          Age      Gendermale  Age:Gendermale
##      28.69400         0.02296        -0.69234        -0.00529
```

Lab 5: Model Selection for Explanatory Models

Scatterplot Matrices

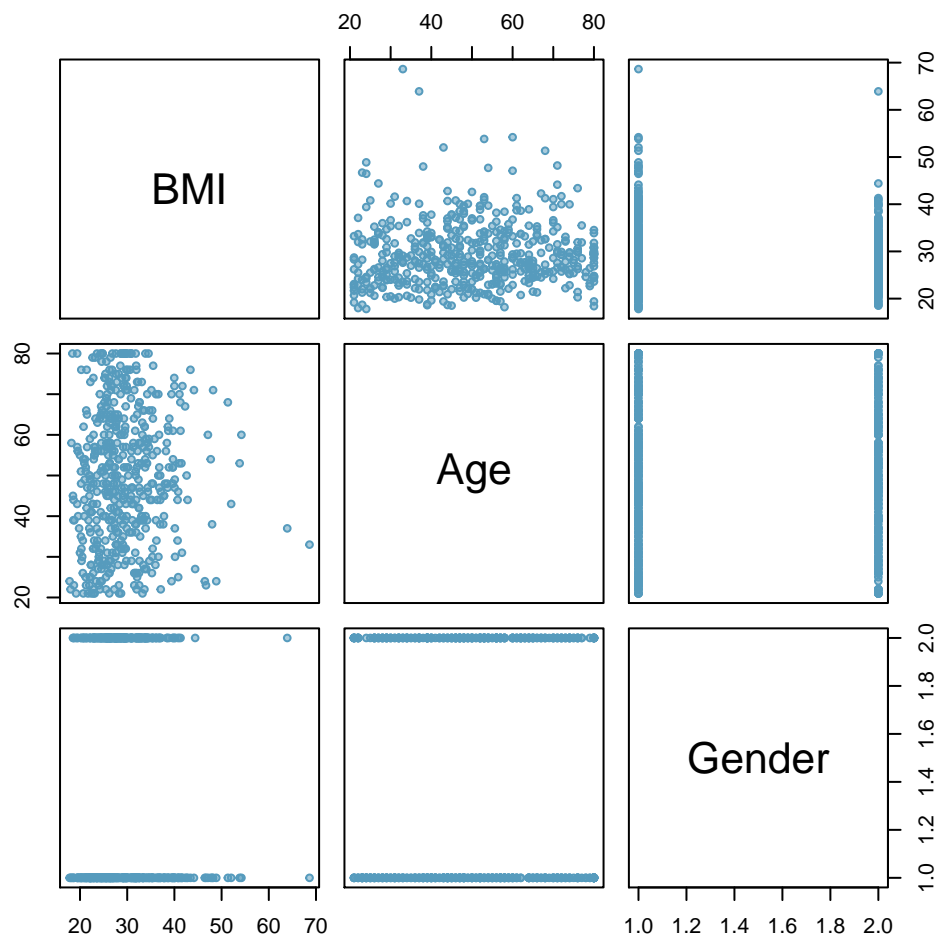
The `pairs()` function creates a matrix of scatterplots. It has the following generic structure

```
pairs(formula, data = )
```

where formula is in the form $\sim x + y + z$, with each term giving a separate variable in the pairs plot and data specifies the data frame for the variables. The function can also take the same graphical parameters that the `plot()` function uses.

The following example shows a scatterplot matrix for the variables BMI, age, and gender in `nhanes.samp.adult.500`.

```
pairs(~ BMI + Age + Gender, data = nhanes.samp.adult.500,  
      pch = 21, cex = 0.7, bg = COL[1, 3], col = COL[1])
```



Correlation Matrices

The `cor()` function was previously introduced to calculate the correlation between two variables. When applied to a matrix, the function will output a correlation matrix of all pairwise correlations.

The `subset()` command has an argument, `select`, useful for pulling out the variables of interest prior to generating a correlation matrix.

In the following examples, the variables BMI, height, and weight are subsetted from `nhanes.samp.adult.500` and stored as a matrix from which a correlation matrix is generated.

```
#subset nhanes.samp.adult.500
nhanes.subset = subset(nhanes.samp.adult.500, select = c(BMI, Height, Weight))

#generate corelation matrix
cor(nhanes.subset, use = "complete.obs")
```

```
##           BMI   Height  Weight
## BMI      1.00000 -0.08178 0.86797
## Height -0.08178  1.00000 0.41023
## Weight  0.86797  0.41023 1.00000
```

Collapsing Factor Levels

The `factor()` function can also be used to collapse levels of a factor.

The following example shows the re-defining of the levels of `grazing.intensity`; the variable initially has five levels (light, less than average, average, moderately heavy, heavy). The first four levels can be combined into a single level `NotHeavy`, while the level `heavy` is renamed `Heavy`.

```
#load data
data("forest.birds")

#view levels of grazing.intensity
levels(forest.birds$grazing.intensity)

## [1] "light"          "less than average" "average"
## [4] "moderately heavy" "heavy"

#create the grazing.binary variable
forest.birds$grazing.binary = forest.birds$grazing.intensity

#redefine the levels of grazing.binary
levels(forest.birds$grazing.binary) = list("NotHeavy" = c("light", "less than average",
                                                         "average", "moderately heavy"),
                                           "Heavy" = "heavy")

#view levels of grazing.binary
levels(forest.birds$grazing.binary)

## [1] "NotHeavy" "Heavy"

#compare tables
table(forest.birds$grazing.intensity)

##
```

```
##          light less than average          average moderately heavy
##          13              8              15              7
##          heavy
##          13
```

```
table(forest.birds$grazing.binary)
```

```
##
## NotHeavy   Heavy
##         43      13
```